A Tractable Model for High-Dimensional Real Sequences

Anonymous Author(s)
Affiliation
Address
`email`

Abstract

We present a probabilistic model for generating and modelling high dimensional real-valued sequences. This model combines a powerful distribution estimator, the Real-Valued Neural Autoregressive Density Estimator (RNADE) with a Recurrent Neural Network (RNN) for capturing temporal dependencies in sequences of high dimensional data. Maximum likelihood learning can be applied to efficiently train the model using a gradient based optimiser. Unlike other models in this family, log-likelihoods for sequences can be computed exactly and efficiently. We evaluate the model's performance on standard datasets and compare its performance to other models.

1 Introduction

Modeling sequences is a fundamental problem in machine learning. Speech, music, video and many other kinds of naturally occurring data are sequential in nature. A good sequential model can be used for discriminative tasks, for sequence completion, as a prior in more complex tasks, sequence denoising and various other applications. There are many existing models such as Linear Dynamical Systems and Hidden Markov Models (HMMs) that model sequential data. Although these models are used extensively in many fields such as speech, audio and music, they can only model a limited a history. Modelling long term dependencies between data points becomes intractable when using these models [10].

Recurrent Neural Networks (RNNs) are simple and powerful models for sequential data. RNNs have an internal memory that allows them to model dependencies between observations separated by a variable number of time steps. In principle RNNs can describe relationships between inputs separated by arbitrary lengths of time. However in practice, training RNNs using gradient based optimisers is not easy and their use in practical applications has been limited. After being ignored by machine learning researchers for a long time, there has been a resurgence of interest in RNNs over the last decade. There have been several key developments in understanding and solving some of the issues associated with gradient based training of RNNs [7, 2]. These advances have made it possible to train RNNs on various tasks and RNNs have been shown to be very successful on a variety of music and language applications [8, 5, 2].

An RNN can either be used to map an input sequence to an output sequence or it can be used as a generative model. Generative models try to model the entire space of inputs given a finite number of training examples. In its simplest form, the RNN can predict unimodal outputs under the assumption that all the output variables are conditionally independent given the remaining outputs. However for many problems this is not satisfactory. A more powerful RNN-based generative model can be constructed by letting the RNN predict the parameters of a powerful distribution estimator. Employing this architecture we obtain complex high-dimensional distributions at each time-step conditioned on the previous inputs. This idea was first employed in the Recurrent Temporal Restricted Boltzmann
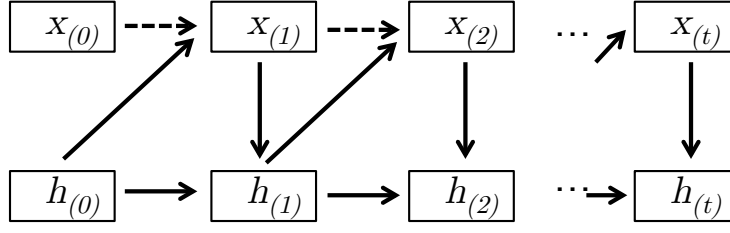
Figure 1: Graphical structure of the generative RNN. Broken arrows indicate optional connections for temporal smoothing.

Machine (RTRBM) [11] model. The model was then extended by combining an RNN with a Neural Autoregressive Distribution Estimator (NADE) and the RBM via a more general architecture [5]. The RNN-RBM and the RNN-NADE have been successfully used for tasks in speech and music [6, 5].

Although the models described above have been successful on various tasks, they are unsuitable for modelling real-valued data. Both the RBM and the NADE are designed to explicitly model binary vectors. If the input data lies in $[0, 1]$, then the RBM can be used via the mean-field approximation. Unbounded data can be modelled by using gaussian visible units and bernoulli hidden units. However these extensions of the RBM to real data have several known limitations [12]. In this paper, we present a new model for real-valued sequential data by combining the RNN with the RNADE. The RNADE is a density estimator that has been shown to outperform gaussian mixture models (GMMs) and a host of other density estimators on several tasks [13]. We show that obtaining probabilities and log-likelihoods from the combined RNN-RNADE model is tractable and fast. We show that the gradients with respect to the model parameters can be calculated exactly and that training can be easily performed with a gradient based optimiser. We evaluate the model's performance on real-valued datasets and demonstrate a marked improvement in performance.

2 Recurrent Neural Networks as generative models

An RNN defines a distribution over an output sequence $\mathbf{x} \equiv \{\mathbf{x}^t \in \mathbb{R}^n, t \leq T\}$ in the following manner:

$$P(\mathbf{x}) = \prod_{t=1}^{T} P(\mathbf{x}^t | \mathcal{A}^t)$$

where $\mathcal{A}^t \equiv \{\mathbf{x}^\tau | \tau < t\}$ is the history of the input sequence $\mathbf{x}$ until time t, $P(\mathbf{x}^t | \mathcal{A}^t)$ is the probability of observing $\mathbf{x}^t$ conditioned on the history of the sequence $\mathcal{A}^t$. RNN based generative models with different properties and capacities can be constructed by carefully choosing the form and parameterization of the conditional $P(\mathbf{x}^t | \mathcal{A}^t)$ as discussed later in this section.

In an RNN with a single hidden layer, the hidden state at time t is given by:

$$\mathbf{h}^t = \sigma(W_{in}\mathbf{x}^t + W_{rec}\mathbf{h}^{t-1} + \mathbf{b}_h) \tag{1}$$

where W_{in} is the weight matrix from the input to the hidden layer, W_{rec} is the recurrent weight matrix from the previous hidden state to the current state and $\mathbf{b}_h$ is the bias vector for the hidden layer. In a standard RNN, the next time step $\mathbf{x}^{t+1}$ is predicted as:

$$\mathbf{x}^{t+1} = \sigma(W_{out}\mathbf{h}^t + \mathbf{b}_{out})$$

where W_{out} connects the hidden layer to the output layer and $\mathbf{b}_{out}$ is the output bias. As shown in Figure 1, there can be optional connections between the input and the input from the previous time-step for temporal smoothing. However as mentioned earlier, each of the outputs is unimodal and independent of the other outputs. These assumptions are too restrictive for many real-world problems and have given rise to a new family of RNN-based generative models with high-dimensional multi-modal conditional distributions at each time-step. Such a model can be realised by letting the hidden layer of the RNN predict parameters of a distribution estimator just like in a mixture density network [4].

2

In previous work [5], the RNN was used to predict the hidden and visible biases of an RBM and a NADE model to yield complex conditional distributions at each time step. Although these models are very good at modelling sequences of binary vectors, they are unsuitable for modelling real-valued data. It is easy to show that any distribution estimator can be used provided the gradients of the cost function with respect to the parameters of the density estimator are obtainable. The combined model can then be trained by maximising the likelihood of the training set and using back-propagation through time (BPTT) [9]. In order to use the power of this architecture for modelling real data, we present a model that combines the RNN and the RNADE. The RNADE is discussed in detail in the following section and the combined model is described in Section 4.

3 The RNADE

The RNADE is a generalisation of the NADE to model real-valued data. Like the NADE, the RNADE expresses the joint probability of the data as a product of one-dimensional conditional distributions as follows:

$$p(x) = \prod_{d=1}^{D} p(x_d|\mathbf{x}_{<\mathbf{d}}) \text{ with } p(x_d|\mathbf{x}_{<\mathbf{d}}) = p_{\mathcal{M}(x_d|\theta_d)}$$

where $p_{\mathcal{M}}$ is a mixture of Gaussians and $\boldsymbol{x}_{<d}$ is a vector of all the dimensions of the data point $< d$. The RNADE is computationally efficient because of the weight sharing employed in the calculation of the hidden state:

$$\mathbf{a}_d = \boldsymbol{W}_{.,<d}\boldsymbol{x}_d + \mathbf{c}$$
$$\boldsymbol{h}_d = \sigma(\rho_d\mathbf{a}_d)$$

where $\mathbf{c} \in \mathbb{R}^H$ and $\boldsymbol{W} \in \mathbb{R}^{D \times (H-1)}$ are neural network parameters that are shared across all the neural networks and $\sigma(x) = 1/(1 + e^{-x})$ is the sigmoid function. $\boldsymbol{W}_{.,<d}$ represents the first $d - 1$ columns of the shared weight matrix. The term ρ_d is a scaling factor which is also learnt from the data. The scaling factor was introduced in [1] in order to prevent the sigmoid hidden units from saturating. The computation of the activations of the hidden units can be made more efficient by performing the computation as:

$$\mathbf{a}_1 = \mathbf{c}, \quad \mathbf{a}_{d+1} = \mathbf{a}_d + x_d\mathbf{W}_{.,d}$$

Unlike the NADE which models each output as a bernouilli distribution, the outputs of each of the feed forward neural networks of the RNADE are mixtures of Gaussians. Therefore the RNADE comprises of D mixture density networks with tied input-to-hidden weights. Once the hidden units of the RNADE have been computed, they are used to compute the parameters of the GMMs $\boldsymbol{\theta}_d = \{\boldsymbol{\alpha}_d, \boldsymbol{\mu}_d, \boldsymbol{\sigma}_d\}$ at each output, where $\boldsymbol{\alpha}_d$ are the mixing coefficients, $\boldsymbol{\mu}_d$ are the means and $\boldsymbol{\sigma}_d$ are the variances. These parameters are computed as follows:

$$\boldsymbol{\alpha}_d = \text{softmax}(\mathbf{V}_d^{\alpha T}\mathbf{h}_d + \mathbf{b}_d^{\alpha})$$
$$\boldsymbol{\mu}_d = \mathbf{V}_d^{\mu T}\mathbf{h}_d + \mathbf{b}_d^{\mu}$$
$$\boldsymbol{\sigma}_d = \exp(\mathbf{V}_d^{\sigma T}\mathbf{h}_d + \mathbf{b}_d^{\sigma})$$

where $\mathbf{V}_d^{\alpha}, \mathbf{V}_d^{\mu}, \mathbf{V}_d^{\sigma}$ are $H \times K$ matrices , $\mathbf{b}_d^{\alpha}, \mathbf{b}_d^{\mu}, \mathbf{b}_d^{\sigma}$ are vectors of size K and K is the number of components in the GMM. More concisely, the RNADE is parameterised by $\mathbf{V}^{\alpha}, \mathbf{V}^{\mu}, \mathbf{V}^{\sigma}$ which are $D \times H \times K$ matrices and $\mathbf{b}^{\alpha}, \mathbf{b}^{\mu}, \mathbf{b}^{\sigma}$ which are matrices of size $D \times K$. The parameters of the RNADE can be learnt by performing gradient ascent on the log-likelihood given a training set.

4 RNN-RNADE

The RNN-RNADE is a sequence of conditional distributions for each time-step of a sequence $\mathbf{x}$. As shown in Figure 2, the parameters of the conditional distributions at each time step t are a function

3
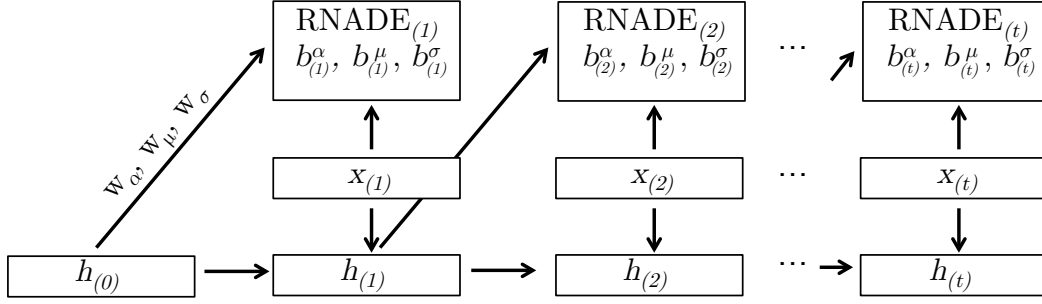
Figure 2: Graphical structure of the RNN-RNADE. The hidden state of the RNN at time t predicts the parameters of the RNADE at $t + 1$.

of the hidden state of the RNN at the previous time-step $t - 1$. We first define the notation in order to simplify discussion of the training algorithm later in the section.

From section 3, the RNADE is parameterised by $\theta \equiv \{\mathbf{V}^\alpha, \mathbf{V}^\mu, \mathbf{V}^\sigma, \mathbf{b}^\alpha, \mathbf{b}^\mu, \mathbf{b}^\sigma\}$. Let $\theta_t \equiv \{\mathbf{V}_t^\alpha, \mathbf{V}_t^\mu, \mathbf{V}_t^\sigma, \mathbf{b}_t^\alpha, \mathbf{b}_t^\mu, \mathbf{b}_t^\sigma\}$ denote the parameters of the conditional RNADE at time t. Although all the parameters of the RNADE at time t can be a function of the hidden state at $t-1$, we consider the matrices $\mathbf{V}^\alpha, \mathbf{V}^\mu, \mathbf{V}^\sigma$ to be fixed at each time step and only consider the case where the biases $\mathbf{b}^\alpha, \mathbf{b}^\mu, \mathbf{b}^\sigma$ to be time-dependent through the RNN hidden state. This constraint significantly reduces the number of parameters that need to be estimated and makes the model and training computationally more efficient. Therefore, $\theta_t \equiv \{\mathbf{V}^\alpha, \mathbf{V}^\mu, \mathbf{V}^\sigma, \mathbf{b}_t^\alpha, \mathbf{b}_t^\mu, \mathbf{b}_t^\sigma\}$. As described later, we experiment with the architecture to ascertain which permutation of time dependent parameters $\mathbf{b}_t^\alpha, \mathbf{b}_t^\mu, \mathbf{b}_t^\sigma$ gives the best results.

The time-dependent RNADE parameters are give by:

$$\text{rvec}(\mathbf{b}_t^\alpha) = \text{rvec}(\mathbf{b}^\alpha) + W_\alpha \mathbf{h}^{t-1} \tag{2}$$

$$\text{rvec}(\mathbf{b}_t^\mu) = \text{rvec}(\mathbf{b}^\mu) + W_\mu \mathbf{h}^{t-1} \tag{3}$$

$$\text{rvec}(\mathbf{b}_t^\sigma) = \text{rvec}(\mathbf{b}^\sigma) + W_\sigma \mathbf{h}^{t-1} \tag{4}$$

where $\text{rvec}(A) = [a_{1,1}, ..., a_{1,n}, a_{2,1}, ..., a_{2,n}, ..., a_{m,1}, ..., a_{m,n}]$ for some $m \times n$ matrix A, $W_\alpha, W_\mu, W_\sigma \in \mathbb{R}^{r \times Dc}$ are matrices from the hidden state of the RNN to the mixing coefficients, means and standard deviations of the RNADE respectively, r is the number of hidden units in the RNN, D is the dimensionality of the input vectors and c is the number of components of each GMM in the RNADE.

4.1 Learning in the RNN-RNADE

The model can be trained by minimising the negative log-likelihood of training sequences using gradient descent. The cost function is given by:

$$L(\theta) = -\log P(\mathbf{x})$$

$$= -\sum_{t=1}^{T} \log P(\mathbf{x}^t; \theta^t) \tag{5}$$

The gradients of the conditionals $P(\mathbf{x}^t; \theta^t)$ with respect to the parameters of the RNADE can be calculated, the details of which are outlined in [13]. Once we obtain the gradients with respect to the time-dependent parameters, the entire network can be trained using BPTT.

From Equation 5, $\frac{\partial L}{\partial \mathbf{b}_t^\alpha} = \frac{\partial(-P(\mathbf{x}^t; \theta^t))}{\partial \mathbf{b}_t^\alpha}$ which can be easily obtained as shown in [13]. The derivatives of the cost with respect to the other time-dependent parameters can be obtained in a similar

4

manner. Once we obtain $\frac{\partial L}{\partial \mathbf{b}_t^\alpha}, \frac{\partial L}{\partial \mathbf{b}_t^\mu}, \frac{\partial L}{\partial \mathbf{b}_t^\sigma}$, the gradients with respect to the other model parameters can be calculated as follows:

$$\frac{\partial L}{\partial W_\alpha} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{b}_t^\alpha} \mathbf{h}^{t-1^T} \tag{6}$$

$$\frac{\partial L}{\partial W_\mu} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{b}_t^\mu} \mathbf{h}^{t-1^T} \tag{7}$$

$$\frac{\partial L}{\partial W_\sigma} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{b}_t^\sigma} \mathbf{h}^{t-1^T} \tag{8}$$

Using equations 1,2,3,4:

$$\frac{\partial L}{\partial \mathbf{h}^t} = W_{rec} \frac{\partial L}{\partial \mathbf{h}^{t+1}} \mathbf{h}^{t+1}(1 - \mathbf{h}^{t+1}) + W_\alpha \frac{\partial L}{\partial \mathbf{b}_{t+1}^\alpha} + W_\mu \frac{\partial L}{\partial \mathbf{b}_{t+1}^\mu} + W_\sigma \frac{\partial L}{\partial \mathbf{b}_{t+1}^\sigma} \tag{9}$$

Using equation 9, the gradients with respect to the remaining model parameters can be calculated easily.

$$\frac{\partial L}{\partial \mathbf{b}_h} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{h}^t} \mathbf{h}^t(1 - \mathbf{h}^t)$$

$$\frac{\partial L}{\partial W_{rec}} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{h}^t} \mathbf{h}^t(1 - \mathbf{h}^t) \mathbf{h}^{t-1^T}$$

$$\frac{\partial L}{\partial W_{in}} = \sum_{t=1}^{T} \frac{\partial L}{\partial \mathbf{h}^t} \mathbf{h}^t(1 - \mathbf{h}^t) \mathbf{x}^{t^T}$$

As shown above, the gradients of the cost function can be calculated and used to train the joint model. In our implementation of the model, we used the automatic differentiation library Theano [3] instead of calculating the gradients by hand. In the supplementary material, we show detailed calculations for the gradient calculations. Unlike the RNN-RBM, the gradients of the RNN-RNADE model can be calculated exactly. Therefore model training can benefit from the use of more powerful gradient based optimiser like the Hessian Free optimiser [7]

5   Experiments

The RNN-RNADE is tested on three tasks, a simple 2d trajectory, videos of bouncing balls and motion capture data. On all three tasks we report the log probabilities as well as the squared prediction error, allowing comparisons between our approach and the RTRBM and RNN-RBM in the latter two tasks. Code for experiments is available in the supplementary material.

5.1   2d stochastic trajectory

This task was used as a demonstrative example of the system. At each time step a particles [x,y] coordinates are determined by the system of equations:

$x = N(sin(t) + 1, 0.1)$
$y = x.N(sin(t + 0.5) + 1, 0.2)$

where t is the timestep. The model is trained on sequences of 100 consecutive timesteps. A model was trained consisting of an RNADE with 1 component per dimension and 20 hidden neurons per dimension, with the RNN employing 20 recurrent neurons. The squared prediction error per frame is 0.16. Videos of predictions of the next time frame, as well as sequence completion can be seen in the supplementary material. A limited grid search was performed to find hyperparameters such as number of hidden and recurrent neurons,

5.2 Motion capture data

As in [11] and [5] we use the motion capture dataset from [11], which consists of sequences of 50 frames, each containing 49 real-valued dimensions representing joint angles, translations and rotations of the base of the spine. Using an RNADE with 1 component per dimension and 200 hidden neurons per dimension and 200 recurrent neurons in the RNN, the RNN-RNADE produces a mean squared test error per frame of 2, substantially lower than 20.1 for RTRBM and 16.2 for the RNN-RBM as reported in [5].

5.3 Videos of bouncing balls

The bouncing balls dataset consists of synthetic videos of three balls bouncing in a box, as described in [11]. We use videos of 15x15 pixels (225 dimensions) and sequences of 128 frames. Each pixel is a real value in the range [0,1]. As the RNADE samples from a Gaussian distribution, it is possible for illegal values to be produced and so the output of the model is constrained to the minimum and maximum values of the data. Using an RNADE with 1 component per dimension and 200 hidden neurons per dimension and 200 recurrent neurons in the RNN, the RNN-RNADE achieves a squared test error per frame of 2.0, similar to the RTRBM but substantially higher than the RNN-RBM as reported in [5]. Videos of errors per frame and generated sequences can be found in the supplementary material.

References

[1] Yoshua Bengio. Discussion of "The Neural Autoregressive Distribution Estimator". In Geoffrey J. Gordon and David B. Dunson, editors, *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS-11)*, volume 15, pages 38–39. Journal of Machine Learning Research - Workshop and Conference Proceedings, 2011.

[2] Yoshua Bengio, Nicolas Boulanger-Lewandowski, and Razvan Pascanu. Advances in optimizing recurrent networks. 2012.

[3] James Bergstra, Olivier Breuleux, Frédéric Bastien, Pascal Lamblin, Razvan Pascanu, Guillaume Desjardins, Joseph Turian, David Warde-Farley, and Yoshua Bengio. Theano: a CPU and GPU math expression compiler. In *Proceedings of the Python for Scientific Computing Conference (SciPy)*, June 2010. Oral Presentation.

[4] Christopher M Bishop. Mixture density networks. 1994.

[5] N. Boulanger-Lewandowski, Y. Bengio, and P. Vincent. Modeling temporal dependencies in high-dimensional sequences: Application to polyphonic music generation and transcription. In *29th International Conference on Machine Learning*, Edinburgh, Scotland, UK, 2012.

[6] Nicolas Boulanger-Lewandowski, Jasha Droppo, Mike Seltzer, and Dong Yu. Phone sequence modeling with recurrent neural networks. In *Proceedings of the 2014 IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP 2014), Florence, Italy*, May 2014.

[7] James Martens and Ilya Sutskever. Learning recurrent neural networks with hessian-free optimization. In *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*, pages 1033–1040, 2011.

[8] Tomas Mikolov, Anoop Deoras, Stefan Kombrink, Lukas Burget, and Jan Cernocky. Empirical evaluation and combination of advanced language modeling techniques. In *INTERSPEECH*, pages 605–608, 2011.

[9] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning internal representations by error propagation. Technical report, DTIC Document, 1985.

[10] Ilya Sutskever and Geoffrey E Hinton. Learning multilevel distributed representations for high-dimensional sequences. In *International Conference on Artificial Intelligence and Statistics*, pages 548–555, 2007.

[11] Ilya Sutskever, Geoffrey E Hinton, and Graham W Taylor. The recurrent temporal restricted boltzmann machine. In *Advances in Neural Information Processing Systems*, pages 1601–1608, 2008.

[12] Lucas Theis, Sebastian Gerwinn, Fabian Sinz, and Matthias Bethge. In all likelihood, deep belief is not enough. *The Journal of Machine Learning Research*, 12:3071–3096, 2011.

[13] Benigno Uria, Iain Murray, and Hugo Larochelle. Rnade: The real-valued neural autoregressive density-estimator. In *Advances in Neural Information Processing Systems 26*, pages 2175–2183. 2013.